

# cnn\_analysis

December 12, 2024

```
[1]: # Install basic modules and make sure they are available with the latest pip
      ↪version.
      # Always updating PIP could be either good or bad, you just have to choose one
      ↪base on the situation around.

      # I use --quiet and --no-warn-script-location to hide the output of my
      ↪directory paths
      import sys
      import os

      ![sys.executable] -m pip install --upgrade pip matplotlib numpy
      ↪tensorflow-macos tensorflow-metal scikit-learn --quiet
      ↪--no-warn-script-location
```

```
[8]: from sklearn.model_selection import train_test_split
      from tensorflow.keras.utils import to_categorical
      import tensorflow as tf
      from tensorflow.keras import layers
      from tensorflow.keras import Model
      import numpy as np
      import os
      from PIL import Image
      import glob
      import matplotlib.pyplot as plt
      from tensorflow.keras import Model
      from tensorflow.keras.optimizers import RMSprop
      from tensorflow.keras.callbacks import LearningRateScheduler, EarlyStopping
      from sklearn.utils.class_weight import compute_class_weight
```

```
[3]: # Define dataset path
      dataset_path = os.path.join(os.getcwd().replace("investigation",
      ↪"kaggledataset"), 'garbage_classification')

      # Load all images and labels
      image_data = []
      labels = []
      class_names = sorted(os.listdir(dataset_path))
```

```

print(f"Classes: {class_names}")

for class_idx, class_name in enumerate(class_names):
    class_folder = os.path.join(dataset_path, class_name)
    if os.path.isdir(class_folder):
        for img_file in glob.glob(os.path.join(class_folder, "*.jpg")):
            try:
                # Open the image, resize, and normalize
                img = Image.open(img_file).convert("RGB").resize((256, 256))
                image_data.append(np.array(img) / 255.0) # Normalize to 0-1
            ↪range
                labels.append(class_idx)
        except Exception as e:
            print(f"Error loading image {img_file}: {e}")

# Convert to NumPy arrays
image_data = np.array(image_data, dtype="float32")
labels = np.array(labels)

# One-hot encode the labels
labels_one_hot = to_categorical(labels, num_classes=len(class_names))

# Split data into 80/20 train/validation
train_data, test_data, train_labels, test_labels = train_test_split(
    image_data, labels_one_hot, test_size=0.2, random_state=42, stratify=labels
)

print(f"Train data shape: {train_data.shape}")
print(f"Train labels shape: {train_labels.shape}")
print(f"Validation data shape: {test_data.shape}")
print(f"Validation labels shape: {test_labels.shape}")

```

```

Classes: ['battery', 'biological', 'brown-glass', 'cardboard', 'clothes',
'green-glass', 'metal', 'paper', 'plastic', 'shoes', 'trash', 'white-glass']
Train data shape: (12412, 256, 256, 3)
Train labels shape: (12412, 12)
Validation data shape: (3103, 256, 256, 3)
Validation labels shape: (3103, 12)

```

```

[4]: num_classes = len(class_names) # Number of classes

# Our input feature map is 150x150x3: 150x150 for the image pixels, and 3 for
# the three color channels: R, G, and B
img_input = layers.Input(shape=(256, 256, 3))

# First convolution extracts 32 filters that are 3x3
# Convolution is followed by max-pooling layer with a 2x2 window

```

```

x = layers.Conv2D(32, 3, activation=None, kernel_regularizer=tf.keras.
↳regularizers.l2(0.01))(img_input)
x = layers.BatchNormalization()(x)
x = layers.ReLU()(x)
x = layers.MaxPooling2D(2)(x)

# Second convolution extracts 64 filters that are 3x3
# Convolution is followed by max-pooling layer with a 2x2 window
x = layers.Conv2D(64, 3, activation=None, kernel_regularizer=tf.keras.
↳regularizers.l2(0.01))(x)
x = layers.BatchNormalization()(x)
x = layers.ReLU()(x)
x = layers.MaxPooling2D(2)(x)

# Third convolution extracts 128 filters that are 3x3
# Convolution is followed by max-pooling layer with a 2x2 window
x = layers.Conv2D(128, 3, activation=None, kernel_regularizer=tf.keras.
↳regularizers.l2(0.01))(x)
x = layers.BatchNormalization()(x)
x = layers.ReLU()(x)
x = layers.MaxPooling2D(2)(x)

# Flatten feature map to a 1-dim tensor so we can add fully connected layers
x = layers.Flatten()(x)
x = layers.Dense(512, activation=None, kernel_regularizer=tf.keras.regularizers.
↳l2(0.01))(x)
x = layers.BatchNormalization()(x)
x = layers.ReLU()(x)
x = layers.Dropout(0.7)(x)

# Create output layer with a single node and sigmoid activation
output = layers.Dense(num_classes, activation='softmax')(x)

```

```

2024-12-12 14:07:03.933364: I metal_plugin/src/device/metal_device.cc:1154]
Metal device set to: Apple M3 Max
2024-12-12 14:07:03.933461: I metal_plugin/src/device/metal_device.cc:296]
systemMemory: 48.00 GB
2024-12-12 14:07:03.933476: I metal_plugin/src/device/metal_device.cc:313]
maxCacheSize: 18.00 GB
2024-12-12 14:07:03.933737: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305]
Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel
may not have been built with NUMA support.
2024-12-12 14:07:03.933762: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271]
Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0
MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id:

```

<undefined>)

```
[5]: # Create model:
model = Model(img_input, output)

model.summary()

# Define optimizer
optimizer = RMSprop(learning_rate=0.0001)

model.compile(
    loss='categorical_crossentropy',
    optimizer=optimizer,
    metrics=['acc']
)
```

Model: "functional"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 256, 256, 3)	0
conv2d (Conv2D)	(None, 254, 254, 32)	896
batch_normalization (BatchNormalization)	(None, 254, 254, 32)	128
re_lu (ReLU)	(None, 254, 254, 32)	0
max_pooling2d (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_1 (Conv2D)	(None, 125, 125, 64)	18,496
batch_normalization_1 (BatchNormalization)	(None, 125, 125, 64)	256
re_lu_1 (ReLU)	(None, 125, 125, 64)	0
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_2 (Conv2D)	(None, 60, 60, 128)	73,856
batch_normalization_2 (BatchNormalization)	(None, 60, 60, 128)	512
re_lu_2 (ReLU)	(None, 60, 60, 128)	0

max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 128)	0
flatten (Flatten)	(None, 115200)	0
dense (Dense)	(None, 512)	58,982,912
batch_normalization_3 (BatchNormalization)	(None, 512)	2,048
re_lu_3 (ReLU)	(None, 512)	0
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 12)	6,156

Total params: 59,085,260 (225.39 MB)

Trainable params: 59,083,788 (225.39 MB)

Non-trainable params: 1,472 (5.75 KB)

```
[10]: class_weights = compute_class_weight('balanced', classes=np.unique(labels),
      ↪y=labels)
class_weights_dict = dict(enumerate(class_weights))

early_stopping = EarlyStopping(monitor='val_acc', patience=20, verbose=1,
      ↪restore_best_weights=True)

def smooth_lr(epoch):
    base_lr = 0.0001
    decay = 0.9 # Slight decay every epoch
    return base_lr * (decay ** epoch)

lr_scheduler = LearningRateScheduler(smooth_lr)

# Cyclical Learning Rate
def clr(epoch):
    base_lr = 0.0001
    max_lr = 0.001
    step_size = 10
    cycle = np.floor(1 + epoch / (2 * step_size))
    x = np.abs(epoch / step_size - 2 * cycle + 1)
    lr = base_lr + (max_lr - base_lr) * max(0, (1 - x))
    return lr
```

```

clr_callback = LearningRateScheduler(clr)

# Class weights
class_weights = compute_class_weight('balanced', classes=np.unique(labels),
    ↪y=labels)
class_weights_dict = dict(enumerate(class_weights))

validation_data = tf.data.Dataset.from_tensor_slices((test_data, test_labels)).
    ↪shuffle(1000).batch(32)

```

```

[11]: result = model.fit(
    train_data,
    train_labels,
    epochs=50,
    batch_size=32,
    validation_data=validation_data,
    verbose=1,
    class_weight=class_weights_dict,
    callbacks=[early_stopping, lr_scheduler]
)

```

Epoch 1/120

2024-12-12 14:09:56.575588: I  
tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:117]  
Plugin optimizer for device\_type GPU is enabled.

388/388                    48s 116ms/step -  
acc: 0.4115 - loss: 11.7814 - val\_acc: 0.3039 - val\_loss: 6.5655 -  
learning\_rate: 1.0000e-04

Epoch 2/120

388/388                    45s 115ms/step -  
acc: 0.6012 - loss: 5.1413 - val\_acc: 0.6262 - val\_loss: 3.7608 - learning\_rate:  
9.0000e-05

Epoch 3/120

388/388                    45s 115ms/step -  
acc: 0.6557 - loss: 3.5513 - val\_acc: 0.7061 - val\_loss: 3.1458 - learning\_rate:  
8.1000e-05

Epoch 4/120

388/388                    45s 115ms/step -  
acc: 0.7098 - loss: 3.0106 - val\_acc: 0.7151 - val\_loss: 2.9199 - learning\_rate:  
7.2900e-05

Epoch 5/120

388/388                    45s 116ms/step -  
acc: 0.7545 - loss: 2.7232 - val\_acc: 0.7048 - val\_loss: 2.8962 - learning\_rate:  
6.5610e-05

Epoch 6/120

388/388                    45s 116ms/step -

acc: 0.7950 - loss: 2.5181 - val\_acc: 0.7093 - val\_loss: 2.7528 - learning\_rate: 5.9049e-05  
Epoch 7/120  
388/388 44s 115ms/step -  
acc: 0.8259 - loss: 2.3311 - val\_acc: 0.7203 - val\_loss: 2.6476 - learning\_rate: 5.3144e-05  
Epoch 8/120  
388/388 45s 115ms/step -  
acc: 0.8573 - loss: 2.1735 - val\_acc: 0.7151 - val\_loss: 2.6300 - learning\_rate: 4.7830e-05  
Epoch 9/120  
388/388 45s 115ms/step -  
acc: 0.8778 - loss: 2.0065 - val\_acc: 0.7396 - val\_loss: 2.4540 - learning\_rate: 4.3047e-05  
Epoch 10/120  
388/388 44s 114ms/step -  
acc: 0.9000 - loss: 1.8737 - val\_acc: 0.6781 - val\_loss: 2.5367 - learning\_rate: 3.8742e-05  
Epoch 11/120  
388/388 44s 114ms/step -  
acc: 0.9145 - loss: 1.7454 - val\_acc: 0.7235 - val\_loss: 2.3802 - learning\_rate: 3.4868e-05  
Epoch 12/120  
388/388 44s 114ms/step -  
acc: 0.9352 - loss: 1.6379 - val\_acc: 0.7045 - val\_loss: 2.4514 - learning\_rate: 3.1381e-05  
Epoch 13/120  
388/388 45s 115ms/step -  
acc: 0.9488 - loss: 1.5374 - val\_acc: 0.7348 - val\_loss: 2.2860 - learning\_rate: 2.8243e-05  
Epoch 14/120  
388/388 45s 115ms/step -  
acc: 0.9582 - loss: 1.4438 - val\_acc: 0.7451 - val\_loss: 2.1651 - learning\_rate: 2.5419e-05  
Epoch 15/120  
388/388 44s 114ms/step -  
acc: 0.9636 - loss: 1.3721 - val\_acc: 0.7103 - val\_loss: 2.3354 - learning\_rate: 2.2877e-05  
Epoch 16/120  
388/388 44s 113ms/step -  
acc: 0.9683 - loss: 1.3024 - val\_acc: 0.7328 - val\_loss: 2.1075 - learning\_rate: 2.0589e-05  
Epoch 17/120  
388/388 44s 113ms/step -  
acc: 0.9773 - loss: 1.2313 - val\_acc: 0.6851 - val\_loss: 2.2281 - learning\_rate: 1.8530e-05  
Epoch 18/120  
388/388 44s 113ms/step -

acc: 0.9813 - loss: 1.1651 - val\_acc: 0.7432 - val\_loss: 2.0240 - learning\_rate:  
 1.6677e-05  
 Epoch 19/120  
 388/388 44s 114ms/step -  
 acc: 0.9799 - loss: 1.1300 - val\_acc: 0.7357 - val\_loss: 2.0494 - learning\_rate:  
 1.5009e-05  
 Epoch 20/120  
 388/388 44s 114ms/step -  
 acc: 0.9876 - loss: 1.0786 - val\_acc: 0.7051 - val\_loss: 2.1905 - learning\_rate:  
 1.3509e-05  
 Epoch 21/120  
 388/388 44s 114ms/step -  
 acc: 0.9882 - loss: 1.0392 - val\_acc: 0.7309 - val\_loss: 1.9612 - learning\_rate:  
 1.2158e-05  
 Epoch 22/120  
 388/388 44s 114ms/step -  
 acc: 0.9910 - loss: 1.0012 - val\_acc: 0.7286 - val\_loss: 2.0012 - learning\_rate:  
 1.0942e-05  
 Epoch 23/120  
 388/388 44s 114ms/step -  
 acc: 0.9920 - loss: 0.9705 - val\_acc: 0.7393 - val\_loss: 1.8963 - learning\_rate:  
 9.8477e-06  
 Epoch 24/120  
 388/388 44s 115ms/step -  
 acc: 0.9935 - loss: 0.9402 - val\_acc: 0.7596 - val\_loss: 1.8224 - learning\_rate:  
 8.8629e-06  
 Epoch 25/120  
 388/388 44s 114ms/step -  
 acc: 0.9962 - loss: 0.9121 - val\_acc: 0.6184 - val\_loss: 2.4651 - learning\_rate:  
 7.9766e-06  
 Epoch 26/120  
 388/388 44s 114ms/step -  
 acc: 0.9960 - loss: 0.8903 - val\_acc: 0.7432 - val\_loss: 1.8721 - learning\_rate:  
 7.1790e-06  
 Epoch 27/120  
 388/388 44s 114ms/step -  
 acc: 0.9958 - loss: 0.8705 - val\_acc: 0.7193 - val\_loss: 1.9879 - learning\_rate:  
 6.4611e-06  
 Epoch 28/120  
 388/388 44s 114ms/step -  
 acc: 0.9957 - loss: 0.8508 - val\_acc: 0.7686 - val\_loss: 1.6975 - learning\_rate:  
 5.8150e-06  
 Epoch 29/120  
 388/388 44s 114ms/step -  
 acc: 0.9975 - loss: 0.8315 - val\_acc: 0.7467 - val\_loss: 1.7885 - learning\_rate:  
 5.2335e-06  
 Epoch 30/120  
 388/388 44s 114ms/step -



acc: 0.9966 - loss: 0.8191 - val\_acc: 0.7580 - val\_loss: 1.6762 - learning\_rate:  
 4.7101e-06  
 Epoch 31/120  
 388/388 44s 114ms/step -  
 acc: 0.9980 - loss: 0.8032 - val\_acc: 0.7544 - val\_loss: 1.7482 - learning\_rate:  
 4.2391e-06  
 Epoch 32/120  
 388/388 44s 114ms/step -  
 acc: 0.9981 - loss: 0.7949 - val\_acc: 0.7751 - val\_loss: 1.6221 - learning\_rate:  
 3.8152e-06  
 Epoch 33/120  
 388/388 44s 113ms/step -  
 acc: 0.9987 - loss: 0.7773 - val\_acc: 0.7686 - val\_loss: 1.6315 - learning\_rate:  
 3.4337e-06  
 Epoch 34/120  
 388/388 44s 114ms/step -  
 acc: 0.9988 - loss: 0.7648 - val\_acc: 0.7667 - val\_loss: 1.6473 - learning\_rate:  
 3.0903e-06  
 Epoch 35/120  
 388/388 44s 113ms/step -  
 acc: 0.9990 - loss: 0.7557 - val\_acc: 0.7705 - val\_loss: 1.6470 - learning\_rate:  
 2.7813e-06  
 Epoch 36/120  
 388/388 44s 113ms/step -  
 acc: 0.9985 - loss: 0.7475 - val\_acc: 0.7686 - val\_loss: 1.6288 - learning\_rate:  
 2.5032e-06  
 Epoch 37/120  
 388/388 44s 114ms/step -  
 acc: 0.9982 - loss: 0.7420 - val\_acc: 0.7757 - val\_loss: 1.6232 - learning\_rate:  
 2.2528e-06  
 Epoch 38/120  
 388/388 44s 114ms/step -  
 acc: 0.9989 - loss: 0.7288 - val\_acc: 0.7699 - val\_loss: 1.6124 - learning\_rate:  
 2.0276e-06  
 Epoch 39/120  
 388/388 44s 114ms/step -  
 acc: 0.9994 - loss: 0.7201 - val\_acc: 0.7728 - val\_loss: 1.5918 - learning\_rate:  
 1.8248e-06  
 Epoch 40/120  
 388/388 44s 114ms/step -  
 acc: 0.9994 - loss: 0.7137 - val\_acc: 0.7641 - val\_loss: 1.6362 - learning\_rate:  
 1.6423e-06  
 Epoch 41/120  
 388/388 44s 114ms/step -  
 acc: 0.9997 - loss: 0.7067 - val\_acc: 0.7657 - val\_loss: 1.6166 - learning\_rate:  
 1.4781e-06  
 Epoch 42/120  
 388/388 44s 114ms/step -

acc: 0.9993 - loss: 0.7024 - val\_acc: 0.7647 - val\_loss: 1.6062 - learning\_rate:  
 1.3303e-06  
 Epoch 43/120  
 388/388 44s 114ms/step -  
 acc: 0.9992 - loss: 0.6956 - val\_acc: 0.7670 - val\_loss: 1.5911 - learning\_rate:  
 1.1973e-06  
 Epoch 44/120  
 388/388 45s 115ms/step -  
 acc: 0.9990 - loss: 0.6928 - val\_acc: 0.7696 - val\_loss: 1.6115 - learning\_rate:  
 1.0775e-06  
 Epoch 45/120  
 388/388 44s 114ms/step -  
 acc: 0.9992 - loss: 0.6863 - val\_acc: 0.7696 - val\_loss: 1.5821 - learning\_rate:  
 9.6977e-07  
 Epoch 46/120  
 388/388 44s 114ms/step -  
 acc: 0.9996 - loss: 0.6829 - val\_acc: 0.7725 - val\_loss: 1.5824 - learning\_rate:  
 8.7280e-07  
 Epoch 47/120  
 388/388 44s 113ms/step -  
 acc: 0.9995 - loss: 0.6773 - val\_acc: 0.7738 - val\_loss: 1.5885 - learning\_rate:  
 7.8552e-07  
 Epoch 48/120  
 388/388 44s 113ms/step -  
 acc: 0.9993 - loss: 0.6772 - val\_acc: 0.7751 - val\_loss: 1.5825 - learning\_rate:  
 7.0697e-07  
 Epoch 49/120  
 388/388 44s 114ms/step -  
 acc: 0.9996 - loss: 0.6719 - val\_acc: 0.7734 - val\_loss: 1.5790 - learning\_rate:  
 6.3627e-07  
 Epoch 50/120  
 388/388 45s 115ms/step -  
 acc: 0.9992 - loss: 0.6696 - val\_acc: 0.7770 - val\_loss: 1.5827 - learning\_rate:  
 5.7264e-07  
 Epoch 51/120  
 388/388 45s 116ms/step -  
 acc: 0.9997 - loss: 0.6654 - val\_acc: 0.7776 - val\_loss: 1.5900 - learning\_rate:  
 5.1538e-07  
 Epoch 52/120  
 388/388 44s 114ms/step -  
 acc: 0.9997 - loss: 0.6643 - val\_acc: 0.7751 - val\_loss: 1.5764 - learning\_rate:  
 4.6384e-07  
 Epoch 53/120  
 388/388 44s 114ms/step -  
 acc: 0.9992 - loss: 0.6628 - val\_acc: 0.7731 - val\_loss: 1.5804 - learning\_rate:  
 4.1746e-07  
 Epoch 54/120  
 388/388 44s 114ms/step -

```

acc: 0.9990 - loss: 0.6615 - val_acc: 0.7731 - val_loss: 1.5827 - learning_rate:
3.7571e-07
Epoch 55/120
388/388          44s 114ms/step -
acc: 0.9994 - loss: 0.6585 - val_acc: 0.7770 - val_loss: 1.5738 - learning_rate:
3.3814e-07
Epoch 56/120
388/388          45s 115ms/step -
acc: 0.9996 - loss: 0.6563 - val_acc: 0.7763 - val_loss: 1.5791 - learning_rate:
3.0433e-07
Epoch 57/120
388/388          44s 114ms/step -
acc: 0.9998 - loss: 0.6538 - val_acc: 0.7757 - val_loss: 1.5752 - learning_rate:
2.7389e-07
Epoch 58/120
388/388          44s 114ms/step -
acc: 0.9997 - loss: 0.6530 - val_acc: 0.7776 - val_loss: 1.5740 - learning_rate:
2.4650e-07
Epoch 59/120
388/388          45s 115ms/step -
acc: 0.9993 - loss: 0.6522 - val_acc: 0.7767 - val_loss: 1.5674 - learning_rate:
2.2185e-07
Epoch 60/120
388/388          44s 113ms/step -
acc: 0.9993 - loss: 0.6554 - val_acc: 0.7767 - val_loss: 1.5693 - learning_rate:
1.9967e-07
Epoch 61/120
388/388          43s 112ms/step -
acc: 0.9997 - loss: 0.6503 - val_acc: 0.7751 - val_loss: 1.5736 - learning_rate:
1.7970e-07
Epoch 62/120
388/388          43s 112ms/step -
acc: 0.9994 - loss: 0.6497 - val_acc: 0.7767 - val_loss: 1.5749 - learning_rate:
1.6173e-07
Epoch 63/120
388/388          43s 112ms/step -
acc: 0.9998 - loss: 0.6476 - val_acc: 0.7767 - val_loss: 1.5748 - learning_rate:
1.4556e-07
Epoch 64/120
388/388          45s 115ms/step -
acc: 0.9998 - loss: 0.6467 - val_acc: 0.7767 - val_loss: 1.5684 - learning_rate:
1.3100e-07
Epoch 65/120
388/388          44s 114ms/step -
acc: 0.9994 - loss: 0.6463 - val_acc: 0.7767 - val_loss: 1.5740 - learning_rate:
1.1790e-07
Epoch 66/120
388/388          44s 114ms/step -

```

```

acc: 0.9991 - loss: 0.6467 - val_acc: 0.7744 - val_loss: 1.5735 - learning_rate:
1.0611e-07
Epoch 67/120
388/388          45s 115ms/step -
acc: 0.9995 - loss: 0.6484 - val_acc: 0.7767 - val_loss: 1.5699 - learning_rate:
9.5500e-08
Epoch 68/120
388/388          44s 114ms/step -
acc: 0.9994 - loss: 0.6462 - val_acc: 0.7741 - val_loss: 1.5728 - learning_rate:
8.5950e-08
Epoch 69/120
388/388          44s 114ms/step -
acc: 0.9995 - loss: 0.6450 - val_acc: 0.7751 - val_loss: 1.5703 - learning_rate:
7.7355e-08
Epoch 70/120
388/388          44s 113ms/step -
acc: 0.9997 - loss: 0.6443 - val_acc: 0.7747 - val_loss: 1.5718 - learning_rate:
6.9620e-08
Epoch 71/120
388/388          44s 113ms/step -
acc: 0.9990 - loss: 0.6466 - val_acc: 0.7738 - val_loss: 1.5714 - learning_rate:
6.2658e-08
Epoch 71: early stopping
Restoring model weights from the end of the best epoch: 51.

```

```

[12]: #####
# Get predictions for the test data
predictions = model.predict(test_data)

# Convert predictions and true labels from one-hot to class indices
predicted_classes = np.argmax(predictions, axis=1)
true_classes = np.argmax(test_labels, axis=1)

# Calculate overall accuracy
overall_accuracy = np.sum(predicted_classes == true_classes) / len(true_classes)
print(f"Overall Test Accuracy: {overall_accuracy:.2f}")


# Calculate per-class accuracy
num_classes = len(class_names)
class_accuracies = []
for class_index in range(num_classes):
    indices = np.where(true_classes == class_index)[0]
    class_correct = np.sum(predicted_classes[indices] == true_classes[indices])
    class_accuracy = class_correct / len(indices) if len(indices) > 0 else 0
    class_accuracies.append(class_accuracy)

# Plot per-class accuracy

```

```
plt.figure(figsize=(10, 6))
plt.bar(class_names, class_accuracies)
plt.title("Per-Class Accuracy")
plt.xlabel("Class")
plt.ylabel("Accuracy")
plt.ylim(0, 1)
plt.xticks(rotation=45)
plt.grid(axis="y")
plt.show()
```

97/97                      2s 17ms/step  
Overall Test Accuracy: 0.78



```
[13]: from sklearn.metrics import classification_report, accuracy_score
import numpy as np

# Get predictions for the test data
predictions = model.predict(test_data)

# Convert predictions and true labels from one-hot encoding to class indices
predicted_classes = np.argmax(predictions, axis=1)
true_classes = np.argmax(test_labels, axis=1)
```

```

# Compute overall accuracy
accuracy = accuracy_score(true_classes, predicted_classes)
print(f"Overall Accuracy: {accuracy:.2f}")

# Compute classification report (includes Precision, Recall, F1-Score)
report = classification_report(true_classes, predicted_classes,
    target_names=class_names)
print("Classification Report:")
print(report)

```

97/97                      2s 16ms/step

Overall Accuracy: 0.78

Classification Report:

	precision	recall	f1-score	support
battery	0.72	0.69	0.70	189
biological	0.75	0.74	0.74	197
brown-glass	0.75	0.70	0.72	122
cardboard	0.75	0.82	0.78	178
clothes	0.92	0.94	0.93	1065
green-glass	0.84	0.86	0.85	126
metal	0.63	0.38	0.47	154
paper	0.79	0.74	0.76	210
plastic	0.56	0.49	0.52	173
shoes	0.69	0.72	0.71	395
trash	0.66	0.78	0.71	139
white-glass	0.58	0.70	0.63	155
accuracy			0.78	3103
macro avg	0.72	0.71	0.71	3103
weighted avg	0.78	0.78	0.77	3103

[ ]: