

# svm\_cnn\_analysis

December 12, 2024

```
[1]: # Install basic modules and make sure they are available with the latest pip
      ↪version.
      # Always updating PIP could be either good or bad, you just have to choose one
      ↪base on the situation around.

      # I use --quiet and --no-warn-script-location to hide the output of my
      ↪directory paths
      import sys
      import os

      ![sys.executable] -m pip install --upgrade pip matplotlib numpy
      ↪tensorflow-macos tensorflow-metal scikit-learn --quiet
      ↪--no-warn-script-location
```

```
[2]: from tensorflow.keras import layers
      from tensorflow.keras import Model
      from tensorflow.keras.optimizers import RMSprop
      from tensorflow.keras.models import Model
      from tensorflow.keras.utils import to_categorical
      from sklearn.model_selection import train_test_split
      from sklearn.svm import SVC
      from sklearn.metrics import accuracy_score, classification_report
      import numpy as np
      import os
      from PIL import Image
      import glob
```

```
[3]: # Load all images and labels
      dataset_path = os.path.join(os.getcwd().replace("investigation",
      ↪"kaggledataset"), 'garbage_classification')

      image_data = []
      labels = []
      class_names = sorted(os.listdir(dataset_path))
      print(f"Classes: {class_names}")

      for class_idx, class_name in enumerate(class_names):
```

```

class_folder = os.path.join(dataset_path, class_name)
if os.path.isdir(class_folder):
    for img_file in glob.glob(os.path.join(class_folder, "*.jpg")):
        try:
            # Open the image, resize, and normalize
            img = Image.open(img_file).convert("RGB").resize((256, 256))
            image_data.append(np.array(img) / 255.0) # Normalize to 0-1
        ↪range
        labels.append(class_idx)
    except Exception as e:
        print(f"Error loading image {img_file}: {e}")

# Convert to NumPy arrays
image_data = np.array(image_data, dtype="float32")
labels = np.array(labels)

# One-hot encode the labels for CNN training
labels_one_hot = to_categorical(labels, num_classes=len(class_names))

# Split data into train/test sets
train_data, test_data, train_labels, test_labels = train_test_split(
    image_data, labels_one_hot, test_size=0.2, random_state=42, stratify=labels
)

```

Classes: ['battery', 'biological', 'brown-glass', 'cardboard', 'clothes', 'green-glass', 'metal', 'paper', 'plastic', 'shoes', 'trash', 'white-glass']

```

[4]: img_input = layers.Input(shape=(256, 256, 3))

# CNN architecture
x = layers.Conv2D(32, 3, activation=None)(img_input)
x = layers.BatchNormalization()(x)
x = layers.ReLU()(x)
x = layers.MaxPooling2D(2)(x)

x = layers.Conv2D(64, 3, activation=None)(x)
x = layers.BatchNormalization()(x)
x = layers.ReLU()(x)
x = layers.MaxPooling2D(2)(x)

x = layers.Conv2D(128, 3, activation=None)(x)
x = layers.BatchNormalization()(x)
x = layers.ReLU()(x)
x = layers.MaxPooling2D(2)(x)

x = layers.Flatten()(x)
x = layers.Dense(512, activation=None)(x)

```

```

x = layers.BatchNormalization()(x)
x = layers.ReLU()(x)
output = layers.Dense(len(class_names), activation='softmax')(x)

cnn_model = Model(img_input, output)

cnn_model.compile(
    loss='categorical_crossentropy',
    optimizer=RMSprop(learning_rate=0.0001),
    metrics=['acc']
)

# Train CNN
cnn_model.fit(
    train_data,
    train_labels,
    epochs=30, # Reduced for demonstration
    batch_size=32,
    validation_data=(test_data, test_labels)
)

```

```

2024-12-12 15:39:42.586247: I metal_plugin/src/device/metal_device.cc:1154]
Metal device set to: Apple M3 Max
2024-12-12 15:39:42.586272: I metal_plugin/src/device/metal_device.cc:296]
systemMemory: 48.00 GB
2024-12-12 15:39:42.586281: I metal_plugin/src/device/metal_device.cc:313]
maxCacheSize: 18.00 GB
2024-12-12 15:39:42.586294: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305]
Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel
may not have been built with NUMA support.
2024-12-12 15:39:42.586305: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271]
Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0
MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id:
<undefined>)

```

Epoch 1/30

```

2024-12-12 15:39:50.303138: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:117]
Plugin optimizer for device_type GPU is enabled.

```

```

388/388          44s 108ms/step -
acc: 0.6094 - loss: 1.2758 - val_acc: 0.5450 - val_loss: 1.4117

```

Epoch 2/30

```

388/388          40s 102ms/step -
acc: 0.8611 - loss: 0.4477 - val_acc: 0.7390 - val_loss: 0.8195

```

Epoch 3/30

388/388                    40s 104ms/step -  
 acc: 0.9510 - loss: 0.1778 - val\_acc: 0.7396 - val\_loss: 0.8717  
 Epoch 4/30  
 388/388                    40s 102ms/step -  
 acc: 0.9841 - loss: 0.0731 - val\_acc: 0.7325 - val\_loss: 0.9369  
 Epoch 5/30  
 388/388                    39s 102ms/step -  
 acc: 0.9913 - loss: 0.0417 - val\_acc: 0.7535 - val\_loss: 0.9677  
 Epoch 6/30  
 388/388                    40s 102ms/step -  
 acc: 0.9966 - loss: 0.0240 - val\_acc: 0.7042 - val\_loss: 1.2862  
 Epoch 7/30  
 388/388                    40s 102ms/step -  
 acc: 0.9970 - loss: 0.0161 - val\_acc: 0.7580 - val\_loss: 0.9367  
 Epoch 8/30  
 388/388                    39s 101ms/step -  
 acc: 0.9993 - loss: 0.0075 - val\_acc: 0.7470 - val\_loss: 1.0077  
 Epoch 9/30  
 388/388                    40s 102ms/step -  
 acc: 0.9986 - loss: 0.0085 - val\_acc: 0.7564 - val\_loss: 0.9552  
 Epoch 10/30  
 388/388                    39s 100ms/step -  
 acc: 0.9996 - loss: 0.0056 - val\_acc: 0.7341 - val\_loss: 1.0580  
 Epoch 11/30  
 388/388                    39s 100ms/step -  
 acc: 0.9993 - loss: 0.0046 - val\_acc: 0.7609 - val\_loss: 0.9937  
 Epoch 12/30  
 388/388                    39s 101ms/step -  
 acc: 0.9995 - loss: 0.0040 - val\_acc: 0.7441 - val\_loss: 1.1035  
 Epoch 13/30  
 388/388                    40s 102ms/step -  
 acc: 0.9994 - loss: 0.0040 - val\_acc: 0.7473 - val\_loss: 1.1447  
 Epoch 14/30  
 388/388                    40s 102ms/step -  
 acc: 0.9990 - loss: 0.0049 - val\_acc: 0.7631 - val\_loss: 0.9834  
 Epoch 15/30  
 388/388                    40s 103ms/step -  
 acc: 0.9995 - loss: 0.0037 - val\_acc: 0.7651 - val\_loss: 1.0743  
 Epoch 16/30  
 388/388                    40s 102ms/step -  
 acc: 0.9997 - loss: 0.0026 - val\_acc: 0.7618 - val\_loss: 1.0169  
 Epoch 17/30  
 388/388                    39s 101ms/step -  
 acc: 0.9997 - loss: 0.0025 - val\_acc: 0.7586 - val\_loss: 1.1095  
 Epoch 18/30  
 388/388                    39s 101ms/step -  
 acc: 0.9992 - loss: 0.0039 - val\_acc: 0.7538 - val\_loss: 1.1540  
 Epoch 19/30

```

388/388          40s 102ms/step -
acc: 0.9998 - loss: 0.0018 - val_acc: 0.7058 - val_loss: 1.3810
Epoch 20/30
388/388          39s 101ms/step -
acc: 0.9994 - loss: 0.0028 - val_acc: 0.7596 - val_loss: 1.0664
Epoch 21/30
388/388          39s 101ms/step -
acc: 0.9996 - loss: 0.0021 - val_acc: 0.7612 - val_loss: 1.0729
Epoch 22/30
388/388          39s 101ms/step -
acc: 0.9998 - loss: 0.0012 - val_acc: 0.7618 - val_loss: 1.0754
Epoch 23/30
388/388          39s 101ms/step -
acc: 0.9994 - loss: 0.0032 - val_acc: 0.7580 - val_loss: 1.0902
Epoch 24/30
388/388          39s 102ms/step -
acc: 0.9982 - loss: 0.0075 - val_acc: 0.7689 - val_loss: 1.0264
Epoch 25/30
388/388          40s 102ms/step -
acc: 0.9997 - loss: 0.0015 - val_acc: 0.7651 - val_loss: 1.1445
Epoch 26/30
388/388          39s 101ms/step -
acc: 0.9992 - loss: 0.0039 - val_acc: 0.7612 - val_loss: 1.0925
Epoch 27/30
388/388          39s 100ms/step -
acc: 0.9993 - loss: 0.0031 - val_acc: 0.7583 - val_loss: 1.1288
Epoch 28/30
388/388          39s 100ms/step -
acc: 0.9999 - loss: 0.0011 - val_acc: 0.7573 - val_loss: 1.2308
Epoch 29/30
388/388          39s 101ms/step -
acc: 0.9997 - loss: 0.0012 - val_acc: 0.7506 - val_loss: 1.1876
Epoch 30/30
388/388          39s 101ms/step -
acc: 0.9994 - loss: 0.0040 - val_acc: 0.7676 - val_loss: 1.1025

```

[4]: <keras.src.callbacks.history.History at 0x36a351a50>

```

[ ]: # Extract features using CNN
feature_extractor = Model(inputs=cnn_model.input, outputs=cnn_model.layers[-2].
    ↪output)

train_features = feature_extractor.predict(train_data)
test_features = feature_extractor.predict(test_data)

# Train an SVM on the extracted features
svm_model = SVC(kernel='linear', class_weight='balanced', probability=True)

```

```

svm_model.fit(train_features, np.argmax(train_labels, axis=1))

# Evaluate
y_pred = svm_model.predict(test_features)
print(f"Accuracy: {accuracy_score(np.argmax(test_labels, axis=1), y_pred)}")
print(classification_report(np.argmax(test_labels, axis=1), y_pred,
    ↳target_names=class_names))

```

```

388/388          9s 22ms/step
97/97           2s 16ms/step
Accuracy: 0.7757009345794392

```

	precision	recall	f1-score	support
battery	0.76	0.69	0.72	189
biological	0.74	0.73	0.73	197
brown-glass	0.75	0.66	0.70	122
cardboard	0.78	0.78	0.78	178
clothes	0.87	0.96	0.91	1065
green-glass	0.84	0.86	0.85	126
metal	0.64	0.49	0.56	154
paper	0.80	0.77	0.78	210
plastic	0.55	0.47	0.51	173
shoes	0.72	0.70	0.71	395
trash	0.69	0.73	0.71	139
white-glass	0.59	0.59	0.59	155
accuracy			0.78	3103
macro avg	0.73	0.70	0.71	3103
weighted avg	0.77	0.78	0.77	3103

```

-----
AttributeError                                Traceback (most recent call last)
Cell In[7], line 16
     13 print(f"Accuracy: {accuracy_score(np.argmax(test_labels, axis=1),
    ↳y_pred)}")
     14 print(classification_report(np.argmax(test_labels, axis=1), y_pred,
    ↳target_names=class_names))
--> 16 svm_model.summary()

AttributeError: 'SVC' object has no attribute 'summary'

```

```

[8]: import matplotlib.pyplot as plt
import numpy as np

# Use feature extractor to get features from the test data

```

```

test_features = feature_extractor.predict(test_data)

# Get predictions for the test features
predicted_classes = svm_model.predict(test_features)

# Convert true labels from one-hot encoding to class indices
true_classes = np.argmax(test_labels, axis=1)


# Calculate overall accuracy
overall_accuracy = np.sum(predicted_classes == true_classes) / len(true_classes)
print(f"Overall Test Accuracy: {overall_accuracy:.2f}")

# Calculate per-class accuracy
num_classes = len(class_names)
class_accuracies = []
for class_index in range(num_classes):
    indices = np.where(true_classes == class_index)[0]
    class_correct = np.sum(predicted_classes[indices] == true_classes[indices])
    class_accuracy = class_correct / len(indices) if len(indices) > 0 else 0
    class_accuracies.append(class_accuracy)

# Plot per-class accuracy
plt.figure(figsize=(10, 6))
plt.bar(class_names, class_accuracies)
plt.title("Per-Class Accuracy")
plt.xlabel("Class")
plt.ylabel("Accuracy")
plt.ylim(0, 1)
plt.xticks(rotation=45)
plt.grid(axis="y")
plt.show()

```

97/97                      2s 15ms/step  
Overall Test Accuracy: 0.78



[ ]: